

Project 2 - Classification Problem

Two of the most widely statistic I methods for analyzing categorical outcome variables linear discriminant analysis and logistic regression. For this project, I selected Linear discriminant analysis and logistic regression. Also, random sampling and stratified K-fold sampling is used for both the model. Here, Accuracy of both the models using both sampling techniques are compared. Accuracy is the proportion of correct predictions over total predictions.

Python code:

```
"""
```

Linear Discriminant Analysis with Training & Test Sampling, Stratified K-Fold Sampling

```
"""
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
from sklearn.model_selection import (train_test_split)
```

```
from sklearn.model_selection import RepeatedStratifiedKFold
```

```
df=pd.read_csv('h2Bank.csv', header=None) # use it when file does not have headers
```

```
# Rename column titles
```

```
df.columns = ['v1', 'v2', 'v3', 'v4', 'v5', 'decision']
```

```
print (df.head()) # see first six rows to check everything
```

```
# Define independent variables and class variables
```

```
X = df[['v1', 'v2', 'v3', 'v4', 'v5']]
```

```
y = df['decision']
```

```
# split dataset into training and testing 70-30 ratio
```

```
X_train, X_test, y_train, y_test=train_test_split(X,y, test_size=0.3)# add fourth parameter  
random_state=10 for seeded random number generation
```

```
print('size of test dataset:',len(X_test), ' size of training dataset: ', len(X_train))
```

```
#Fit the LDA model
```

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
model = LinearDiscriminantAnalysis()
```

```
model.fit(X_train, y_train) #learn Discriminant Function
```

```
# Print Training and Test Accuracies
```

```
result1 = model.score(X_train, y_train)
```

```
print(("LDA Training Accuracy: %.2f%%") % (result1*100))
```

```
result = model.score(X_test, y_test)
```

```
print(("LDA Test Accuracy: %.2f%%") % (result*100.0))
```

```
from sklearn.metrics import confusion_matrix
```

```
#Create Training and Test Dataset
```

```
X_train, X_test, y_train, y_test=train_test_split(X,y,random_state=1)
```

```
y_pred=model.predict(X_test) # predict test dataset
```

```
confusion=confusion_matrix(y_test,y_pred)#,labels=[0,1])
```

```
print(confusion) # Column is Actual and Row Title is Predicted
```

```
# Cross Validation using 10 fold sampling
```

```
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
```

```
model1 = LinearDiscriminantAnalysis()
```

```
model1.fit(X, y)
```

```
#evaluate model

from sklearn.model_selection import cross_val_score

scores = cross_val_score(model1, X, y, scoring='accuracy', cv=cv, n_jobs=-1)

print(("Stratified KFold Accuracy: %.2f%%" % (np.mean(scores)*100))
```

Output:

```

v1  v2  v3  v4  v5 decision
0 2.463 0.056 1.682 -0.108 1.173    1
1 1.234 0.091 2.107 0.046 0.144    1
2 2.852 0.056 1.782 0.107 0.668    1
3 -0.775 -0.043 1.372 0.026 0.543    1
4 1.573 0.047 1.377 0.177 1.797    1

size of test dataset: 30 size of training dataset: 70

LDA Training Accuracy: 65.71%

LDA Test Accuracy: 60.00%

[[7 8]

 [7 3]]

Stratified KFold Accuracy: 65.67%
```

"""

Logistic Regression with Training & Test Sampling, Stratified KFold Sampling

"""

```
import pandas as pd

import numpy as np

from sklearn.model_selection import (train_test_split)

from sklearn.model_selection import RepeatedStratifiedKFold


df=pd.read_csv('h2Bank.csv', header=None) # use it when file does not have headers

# Rename column titles
```

```

df.columns = ['v1', 'v2', 'v3', 'v4', 'v5', 'decision']

print (df.head()) # see first six rows to check everything


# Define independent variables and class variables
X = df[['v1', 'v2', 'v3', 'v4', 'v5']]
y = df['decision']


# split dataset into training and testing 70-30 ratio
X_train, X_test, y_train, y_test=train_test_split (X,y, test_size=0.3)# add fourth parameter
random_state=10 for seeded random number generation

print('size of test dataset:',len(X_test), ' size of training dataset: ', len(X_train))


# Logistic Regression with Area Under the Curve-may not work with Iris due to dataset loading
from sklearn.linear_model import LogisticRegression

clf = LogisticRegression(solver="liblinear", random_state=0).fit(X_train, y_train)

from sklearn.metrics import roc_auc_score

print("ROC-AUC score: %.3f" % roc_auc_score(y, clf.predict_proba(X)[:, 1]))


# Print Training and Test Accuracies
result1 = clf.score(X_train, y_train)

print(("LR Training Accuracy: %.2f%%" % (result1*100))

result = clf.score(X_test, y_test)

print(("LR Test Accuracy: %.2f%%" % (result*100))


# Cross Validation using 10 fold sampling
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)


model = LogisticRegression(solver="liblinear", random_state=0)

model.fit(X, y)

```

```
#evaluate model

from sklearn.model_selection import cross_val_score

scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)

print(("Stratified KFold Accuracy: %.2f%%" % (np.mean(scores)*100))
```

Output:

```

v1  v2  v3  v4  v5 decision
0 2.463 0.056 1.682 -0.108 1.173    1
1 1.234 0.091 2.107 0.046 0.144    1
2 2.852 0.056 1.782 0.107 0.668    1
3 -0.775 -0.043 1.372 0.026 0.543    1
4 1.573 0.047 1.377 0.177 1.797    1

size of test dataset: 30 size of training dataset: 70

ROC-AUC score: 0.746

LR Training Accuracy: 75.71%

LR Test Accuracy: 63.33%

Stratified KFold Accuracy: 67.00%
```

Overall Result:

Accuracy %	Random Sampling	Stratified K-fold Sampling
Linear Discriminant Analysis	60.00%	65.67%
Logistic Regression	63.33%	67.00%

Comparing the results of both linear discriminant analysis and logistic regression, it is evident that logistic regression using stratified K-fold sampling yields better accuracy compared to linear discriminant analysis. While both are appropriate for the development of linear classification models, linear discriminant analysis makes more assumptions about the underlying data. Hence, it is assumed that logistic regression is the more flexible and more robust method. Hence, I consider Logistic regression is better compared to LDA. Similarly, stratified Kfold sampling is better compared to random sampling.

Screenshots:

The screenshot shows the Spyder Python IDE with a Python 3.8 environment. The code in the editor performs LDA on a dataset with five features (v1, v2, v3, v4, v5) and a decision variable. The results are displayed in the Variable explorer and the console.

Code (HW2_LDA.py):

```
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import numpy as np
7 from sklearn.model_selection import (train_test_split)
8 from sklearn.model_selection import RepeatedStratifiedKFold
9
10 df=pd.read_csv('h2Bank.csv', header=None) # use it when file does not have headers
11
12 # Rename column titles
13 df.columns = ['v1', 'v2', 'v3', 'v4', 'v5', 'decision']
14 print(df.head()) # see first six rows to check everything
15
16 # Define independent variables and class variables
17 X = df[['v1', 'v2', 'v3', 'v4', 'v5']]
18 y = df['decision']
19
20 # split dataset into training and testing 70-30 ratio
21 X_train, X_test, y_train, y_test=train_test_split(X,y, test_size=0.3) # add fourth parameter random
22 print('size of test dataset:',len(X_test), ' size of training dataset: ', len(X_train))
23
24 #Fit the LDA model
25 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
26 model = LinearDiscriminantAnalysis()
27 model.fit(X_train, y_train) #Learn Discriminant Function
28
29 # Print Training and Test Accuracies
30 result1 = model.score(X_train, y_train)
31 print(("LDA Training Accuracy: %.2f%%" % (result1*100)))
32
33 result = model.score(X_test, y_test)
34 print(("LDA Test Accuracy: %.2f%%" % (result*100.0)))
35
36 from sklearn.metrics import confusion_matrix
37 #Create Training and Test Dataset
38 X_train, X_test, y_train, y_test=train_test_split(X,y,random_state=1)
39 y_pred=model.predict(X_test) # predict test dataset
40 confusion=confusion_matrix(y_test,y_pred),labels=[0,1])
41 print(confusion) # Column is Actual and Row Title is Predicted
42
43 # Cross Validation using 10 fold sampling
44 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
```

Variable explorer:

Name	Type	Size	Value
confusion	Array of int64	(2, 2)	[[7 8] [7 3]]
cv	model_selection._split.RepeatedStratifiedKFold	1	RepeatedStratifiedKFold object of sk...
df	DataFrame	(100, 6)	Column names: v1, v2, v3, v4, v5, decision
model	discriminant_analysis.LinearDiscriminantAnalysis	1	LinearDiscriminantAnalysis object of sklearn.discriminant_analysis mod ...
model1	discriminant_analysis.LinearDiscriminantAnalysis	1	LinearDiscriminantAnalysis object of sklearn.discriminant_analysis mod ...
result	float64	1	0.6
result1	float64	1	0.6571428571428571
scores	Array of float64	(30,)	[0.9 0.5 0.8 ... 0.6 0.8 0.6]
X	DataFrame	(100, 5)	Column names: v1, v2, v3, v4, v5

Console:

```
In [74]: runfile('C:/Users/marut/Desktop/Fall 2021/INFSY 566/Programs/HW2_LDA.py', wdir='C:/Users/marut/Desktop/Fall 2021/INFSY 566/Programs')
v1 v2 v3 v4 v5 decision
0 2.463 0.056 1.682 -0.108 1.173 1
1 1.234 0.091 2.107 0.046 0.144 1
2 2.852 0.056 1.782 0.107 0.668 1
3 -0.775 -0.043 1.372 0.026 0.543 1
4 1.573 0.047 1.377 0.177 1.797 1
size of test dataset: 30 size of training dataset: 70
LDA Training Accuracy: 65.71%
LDA Test Accuracy: 60.00%
[[7 8]
 [7 3]]
Stratified KFold Accuracy: 65.6%
```

The screenshot shows the Spyder Python IDE with a Python 3.8 environment. The code in the editor performs Logistic Regression on the same dataset as the LDA model. The results are displayed in the Variable explorer and the console.

Code (HW2_LR.py):

```
4 import pandas as pd
5 import numpy as np
6 from sklearn.model_selection import (train_test_split)
7 from sklearn.model_selection import RepeatedStratifiedKFold
8
9 df=pd.read_csv('h2Bank.csv', header=None) # use it when file does not have headers
10
11 # Rename column titles
12 df.columns = ['v1', 'v2', 'v3', 'v4', 'v5', 'decision']
13 print(df.head()) # see first six rows to check everything
14
15 # Define independent variables and class variables
16 X = df[['v1', 'v2', 'v3', 'v4', 'v5']]
17 y = df['decision']
18
19 # split dataset into training and testing 70-30 ratio
20 X_train, X_test, y_train, y_test=train_test_split(X,y, test_size=0.3) # add fourth parameter random
21 print('size of test dataset:',len(X_test), ' size of training dataset: ', len(X_train))
22
23 # Logistic Regression with Area Under the Curve may not work with Iris due to dataset loading
24 from sklearn.linear_model import LogisticRegression
25 clf = LogisticRegression(solver='liblinear', random_state=0).fit(X_train, y_train)
26 from sklearn.metrics import roc_auc_score
27 print("ROC-AUC score: %.3f" % roc_auc_score(y, clf.predict_proba(X)[:, 1]))
28
29 # Print Training and Test Accuracies
30 result1 = clf.score(X_train, y_train)
31 print(("LR Training Accuracy: %.2f%%" % (result1*100)))
32
33 result = clf.score(X_test, y_test)
34 print(("LR Test Accuracy: %.2f%%" % (result*100)))
35
36 # Cross Validation using 10 fold sampling
37 cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
38
39 model = LogisticRegression(solver='liblinear', random_state=0)
40 model.fit(X, y)
41
42 #evaluate model
43 from sklearn.model_selection import cross_val_score
44 scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv, n_jobs=-1)
45 print(("Stratified KFold Accuracy: %.2f%%" % (np.mean(scores)*100)))
```

Variable explorer:

Name	Type	Size	Value
clf	linear_model_logistic.LogisticRegression	1	LogisticRegression object of sklearn.linear_model_logistic module
cv	model_selection._split.RepeatedStratifiedKFold	1	RepeatedStratifiedKFold object of sklearn.model_selection._split modul ...
df	DataFrame	(100, 6)	Column names: v1, v2, v3, v4, v5, decision
model	linear_model_logistic.LogisticRegression	1	LogisticRegression object of sklearn.linear_model_logistic module
result	float64	1	0.6333333333333333
result1	float64	1	0.7571428571428571
scores	Array of float64	(30,)	[0.9 0.5 0.8 ... 0.8 0.7 0.6]
X	DataFrame	(100, 5)	Column names: v1, v2, v3, v4, v5
X_test	DataFrame	(30, 5)	Column names: v1, v2, v3, v4, v5

Console:

```
In [27]: runfile('C:/Users/marut/Desktop/Fall 2021/INFSY 566/Programs/HW2_LR.py', wdir='C:/Users/marut/Desktop/Fall 2021/INFSY 566/Programs')
v1 v2 v3 v4 v5 decision
0 2.463 0.056 1.682 -0.108 1.173 1
1 1.234 0.091 2.107 0.046 0.144 1
2 2.852 0.056 1.782 0.107 0.668 1
3 -0.775 -0.043 1.372 0.026 0.543 1
4 1.573 0.047 1.377 0.177 1.797 1
size of test dataset: 30 size of training dataset: 70
ROC-AUC score: 0.746
LR Training Accuracy: 75.71%
LR Test Accuracy: 63.33%
Stratified KFold Accuracy: 67.00%
```